

# Starcode: an exact algorithm for sequence clustering

Eduard Valera Zorita<sup>1,2</sup>, Pol Cuscó<sup>1,2</sup> and Guillaume Filion<sup>1,2\*</sup>

<sup>1</sup>Genome Architecture, Gene Regulation, Stem Cells and Cancer Programme, Centre for Genomic Regulation (CRG), Dr. Aiguader 88, 08003 Barcelona, Spain.

<sup>2</sup>Universitat Pompeu Fabra (UPF), Barcelona, Spain.

Received on XXXXX; revised on XXXXX; accepted on XXXXX

Associate Editor: XXXXXXXX

## ABSTRACT

**Motivation:** The increasing throughput of sequencing technologies offers new applications and challenges for computational biology. One such application is the use of random barcodes to trace and quantify transcripts or lineages in experimental setups. The high error rate of modern sequencers calls for additional post-processing techniques capable of detecting and reverting the misreads. However, in the absence of a reference population, the problem amounts to performing a pairwise comparison of all the barcodes, which is unfeasible for excessive computational complexity.

**Results:** Here we address this problem and describe an exact algorithm to determine which pairs of sequences lie within a given Levenshtein distance. The matched pairs are merged into clusters represented by a canonical sequence. The speed of starcode is attributable to the poucet search, a novel implementation of the Needleman-Wunsch algorithm performed on the nodes of a trie. Parallelization achieves linear scaling of performance on multi-core machines. On the task of clustering random barcodes, starcode outperforms short read mappers and sequence clustering algorithms in both speed and precision.

**Availability and implementation:** The C source code is available at <http://github.com/guil11aume/starcode>.

**Contact:** guillaume.filion@gmail.com

## 1 INTRODUCTION

Sequence clustering is the process of grouping similar biological sequences. It has been traditionally applied to identify related protein families and to reduce sequence redundancy in databases. Recently, the advent of high throughput sequencing has created additional needs for efficient clustering algorithms, in particular because of the high error rate of such technologies. For instance, the Illumina platform (Margulies *et al.*, 2005) shows a 1-2% error rate consisting of substitutions near the 3' end of the read (Dohm *et al.*, 2008; Nakamura *et al.*, 2011). The PacBio platform shows a 15% error rate that mostly correspond to insertions and deletions (Eid *et al.*, 2009). As a consequence, the same sequence is often decoded in different ways, which artificially increases the diversity of the output.

Sequencing errors can be discovered by mapping the reads onto a reference, if it is available. But when the sequences are random or drawn from an unknown reference, clustering is the best option

to tell real from spurious reads. One such case is the use of random barcodes to track cells or transcripts (Schepers *et al.*, 2008; Akhtar *et al.*, 2013). Sequencing errors will create new spurious barcodes that have to be reverted to the original sequence by clustering. Hands on experience with real datasets produced in our laboratory showed that this step can become limiting when the number of unique sequences is high. In search for a solution to this problem, we realized that heuristic approaches rely on assumptions that will perhaps not hold as technologies evolve, we therefore set out to find an exact algorithm.

The first task of clustering is a matching phase where similar barcodes are paired, similarly to linked nodes on a graph. The second task is the clustering proper, where communities in this graph are merged. We called our algorithm “starcode”, in reference to the star shape of the graph formed by the barcodes in the same cluster. It occurred to us that the first version of our algorithm had the same expected running time when barcodes were sorted or not, because it was not exploiting data structuring of any kind. We reasoned that there should be a way to use the prefix redundancy of sorted barcodes to gain speed. This is how we came up with the poucet search algorithm at the heart of the matching step.

Here we describe the starcode algorithm and we benchmark it against existing software that perform related tasks. We show that starcode is faster than the alternatives, and that as an exact algorithm it makes no mistake on ideal datasets. We show that starcode can also be used more generally to cluster k-mers from large datasets from bacterial and even metazoan genomes.

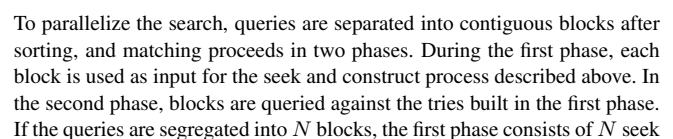
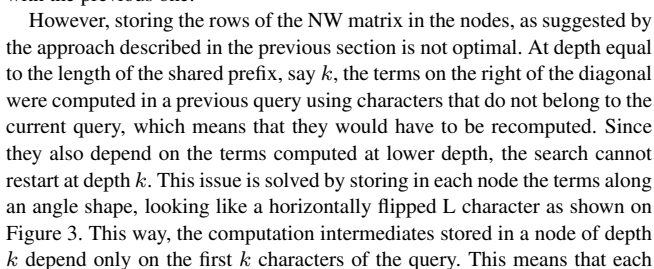
## 2 METHODS


### 2.1 Inexact string matching using tries

The matching method of starcode is a variation of the Needleman-Wunsch (NW) algorithm (Needleman and Wunsch, 1970). In the original algorithm (Figure 1a), the Levenshtein distance between two sequences is found by applying a recurrence relation throughout a matrix of  $mn$  terms, where  $m$  and  $n$  are the respective sequence lengths. The complexity of this dynamic programming approach is  $O(mn)$ .

In many instances, the information of interest is whether the sequences are  $d$ -neighbors (separated by a distance less than or equal to a threshold  $d$ ). In that case, the complexity can be reduced to  $O(d \cdot \min(m, n))$  as described below (Ukkonen, 1995). Instead of initializing the margins and computing all the terms, the matrix is initialized as shown in Figure 1b and only the terms around the diagonal are computed. If the computation of a diagonal term yields a value greater than  $d$ , the process is halted (and the distance

\*to whom correspondence should be addressed





**Fig. 3.** Poucet search algorithm. The algorithm proceeds with the same principles as shown on Figure 2 with the difference that the NW matrix is not updated row-wise, but along horizontally flipped L shapes. As the depth-first search proceeds, these values are stored in the nodes of the trie. Only the nodes at the top contain initialized values; for the other nodes, the values at the border are implicitly known to be 3. Since the values in the vertical part of the flipped L are the same for every child of the same node, they are computed only once (purple arrow). The values in the grey cells will be computed as the search path (red) visits the node. Storing the intermediates in the nodes allows the next query to restart at depth  $k$  if it shares a common prefix of length  $k$  with the current query.

and construct jobs, and the second consists of  $N(N-1)/2$  query jobs. Since the jobs show little dependence on each other, the matching algorithm can be efficiently parallelized provided  $N$  is larger than the number of independent threads.

## 2.5 Clustering

Starcode implements a multi-purpose clustering algorithm called “sphere clustering”, and a message passing algorithm (MacKay, 2002) tailored for the task of clustering random barcodes. In sphere clustering, barcodes are sorted by frequency and each barcode, starting from the most frequent, can claim its  $d$ -neighbors that were not already claimed. This is the method used by Akhtar *et al.* (2013).

In message passing clustering, read counts are distributed equally among the closest neighbors of each barcode if they are at least 5 times more frequent. The sequences with positive counts at the end of the process are considered canonical, and their associated count is the estimated cluster size. A barcode is assigned to a cluster if all its read count are eventually given to the corresponding canonical barcode. The barcodes for which the read counts are split between different canonical barcodes are not assigned to any cluster. The reason for imposing a factor 5 or larger in order to transfer the read counts is that barcodes with similar frequencies are not likely derived from each other through sequencing errors. More likely they are either unrelated, or they both derive from another more abundant barcode.

## 2.6 Benchmark options

The benchmark dataset is described in the Results section. Starcode was run with options `-t 20`. CD-HIT version 4.6.1 was run with options `-c 0.925 -T 20 -M 0`. USEARCH 32 bit version 7.0.1090 was run with options `-threads 20 -cluster_fast -id 0.925`. The Bowtie2 indexer was run as `bowtie2-build --quiet`, and Bowtie2 version 2.1.0 was then run with options `-f -a -p 20 --no-hd --very-sensitive`. The BWA indexer was run as `bwa index -a is` and then BWA version 0.7.9a was run with options `mem -t 20`

`-a -k1 -B0`. The GEM indexer was first run as `gem-indexer -T 20`. GEM version 1.423 (beta) was then run with options `-e3 -s -T 20 --granularity 100000`. The sequential algorithm of Akhtar *et al.* (2013) was implemented in R and relied on the `stringdist` function of the `stringdist` package, used with options `method='lv'`, `maxDist=2`.

## 3 RESULTS

### 3.1 Benchmark against sequence clustering algorithms

Sequence clustering is routinely used to curate databases of non redundant nucleotide or protein sequences. We benchmarked starcode against the two popular sequence clustering algorithms CD-HIT (Fu *et al.*, 2012) and USEARCH (Edgar, 2010).

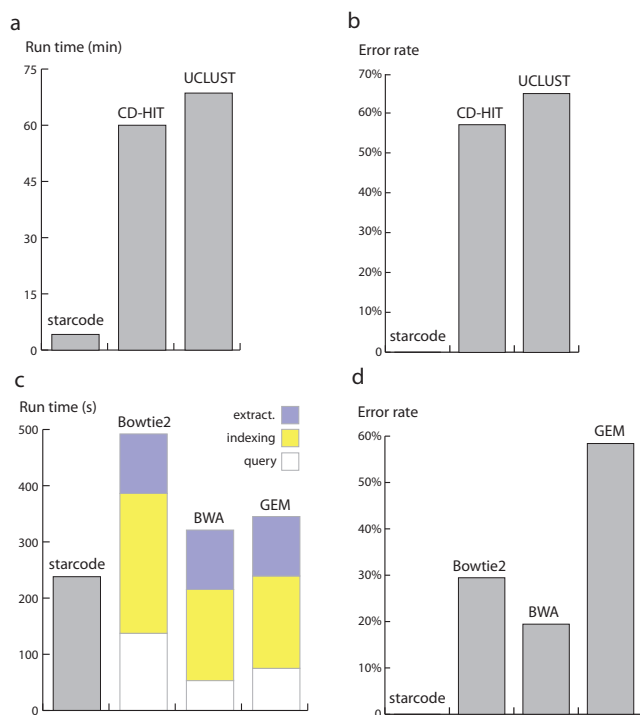
We generated 1 million random 40-mers that were duplicated 47 times. For each 40-mer, we added 3 mutant sequences to the pool by sampling again 3 nucleotides at random positions. The probability that such a mutant sequence is within a distance 3 of another 40-mer is of the order of 10-12 and can be neglected. Each cluster thus consists of 50 sequences, of which 3 have a distance at most 3 from the canonical representative.

Using 20 cores, starcode clustered the shuffled 50 million sequences in 4 minutes without a single error (the output consisted of 1 million clusters of size 50). We submitted the unique sequences of the dataset to CD-HIT and USEARCH. The results are summarized in Figure 4. The running times of CD-HIT and USEARCH were 59 and 64 minutes respectively, and they identified 2,362,352 and 2,739,954 clusters respectively. Both algorithms rely on heuristics, which is why the number of clusters is off target. These results show that starcode outperforms other clustering algorithms in speed and accuracy on this particular task.

### 3.2 Benchmark against short read mappers

A potential strategy to cluster short sequences is to use last generation short read mappers to match a set of sequences against itself. Once matches are available, several community detection algorithms can be used to identify the clusters. We benchmarked starcode against the short read mappers Bowtie2 (Langmead and Salzberg, 2012), BWA (Li and Durbin, 2009) and GEM (Marco-Sola *et al.*, 2012). Because short read mappers do not perform clustering, we only evaluated their performance on the matching problem.

For this test, we used again the dataset described above, in which every sequence has at least 2 distinct matches at a distance 3. This means that we can use matches with a single hit (themselves) as a lower bound on the inaccuracy. We used unique sequences to build an index and queried the index with the same sequences. Figure 4 shows the running time of starcode along with the decomposition of the running time for the three mappers. Note that time to extract unique sequences is always the same because we used the Linux `sort -u` command in every case. It appears that the running time of starcode is about the same as the time required by BWA (the fastest aligner in this test) to index and run the query proper. Figure 4 shows the lower bound on the error rate, as estimated by the proportion of single hits. BWA was found to be the most accurate aligner with a lower bound of 19% errors. This measure is an underestimate of the true error rate, but given that starcode makes no mistake on this dataset, a more accurate measure is not necessary. In conclusion, starcode runs slightly faster than short read mappers



**Fig. 4.** Benchmark of starcode against sequence clustering algorithms and short read mappers. **a** Running time compared to CD-HIT and UCLUST on the dataset described in the text. **b** Error rate compared to CD-HIT and UCLUST, expressed as the percentage of clusters failing to be merged. **c** Running time compared to Bowtie2, BWA and GEM on this dataset. The running time is decomposed in extraction of unique sequences (purple), Burrows-Wheeler indexing (yellow), and query proper (white). The extraction was performed with the Linux sort command, so the time is the same in every case. **d** Error rate compared to Bowtie2, BWA and GEM, expressed as the proportion of matches with a single hit. This is an underestimate of the true error rate.

when all steps are taken into account, and clearly outperforms them in precision on this particular problem.

### 3.3 Clustering TRIP barcodes

In the course of setting up the TRIP technology in our lab (Akhtar *et al.*, 2013), we realized the need to develop efficient algorithms to cluster similar sequences. Briefly, the principle of TRIP (Thousands of Reporters Integrated in Parallel) is to tag reporter transcripts with random barcodes and measure the abundance of barcodes in the RNA as a proxy for gene expression. There is no reference to match aberrant barcodes against, because the tagging sequences are unknown. Instead, barcodes are matched against each other and clustered by similarity to infer canonical sequences.

We tested the efficiency of starcode on the TRIP dataset from Akhtar *et al.* (2013). In the experiment labelled mPGKA, we identified 24.1 million barcode-containing reads out of 26.6 million (91%), consisting of approximately 223,000 and 220,000 unique barcode sequences for replicates 1 and 2 respectively. Following the authors, we kept only the barcodes with at least 5 counts and performed clustering as described: “First we sorted barcodes according to their counts. Then, for each barcode (starting from

the most frequent one), we identified and removed all its mutant versions, defined as barcodes within a Hamming distance of 2.” This method is here on referred to as the “sequential algorithm”. When replacing the Hamming distance by the Levenshtein distance, starcode produced exactly the same output as the sequential algorithm. The running time of starcode averaged over the replicates was 2.90 seconds, instead of more than 3 minutes 40 seconds, which represents a 75 fold speedup.

The performance of the sequential algorithm relies on the arbitrary exclusion of reads with less than 5 counts, the main purpose of which is to reduce the computational burden. When all barcodes were kept, the average running time of starcode was 10.27 seconds, versus > 6 hours for the sequential algorithm. Using the Hamming distance as the authors originally suggested decreased the average running time of the sequential algorithm to 3 hours and 30 minutes (and the output differed from that of starcode). Note that in all the cases mentioned above, starcode was run with a single core to compare the algorithms based on similar computer resources. In conclusion, the barcode clustering problem can be simplified by various heuristics, but starcode brings down the running time to nearly instantaneous, and thereby obviates the need for such arbitrary heuristics.

### 3.4 Identifying enriched sequence motifs

Sequence motifs are thought to play an important role in DNA metabolism. Key regulators such as transcription factors, nucleosomes and non coding RNAs have sequence preferences targeting them to the sites where they act. Identifying those sequences is a way to pinpoint the regulators and the mechanisms they are involved in. However, the sequence motifs are not strictly identical between sites, which is why they are better identified by inexact matching. This problem becomes computationally difficult for long motifs (above 12-13 nucleotides) because of the combinatorial explosion. But as motifs become longer, the problem of identifying abundant inexact matches becomes similar to barcode clustering.

We set up a test based on the meningitis-causing agent *Neisseria meningitidis*. The genome of this bacterium is interspersed with a frequent 12 bp sequence known as DNA uptake sequence (Smith *et al.*, 1999). We extracted the 12-mers from both orientations of the 2.19 Mb genome, yielding 4.39 million 12-mers, consisting of 2.77 million unique sequences. Using starcode to cluster the 12-mers within a Levenshtein distance of 2 took less than 3 minutes with 20 cores. We identified the known DNA uptake sequence of *Neisseria meningitidis* (ATGCCGTCTGAA) as the most abundant 12-mer, comprising 1,466 exact matches and 4,625 inexact matches. This shows that starcode can be used to identify biologically relevant motifs in bacterial genomes.

To test starcode on the scale of a metazoan genome, we used a *Drosophila* dataset from Filion *et al.* (2010). That study describes two signatures of chromatin proteins called Red and Yellow, that both correspond to transcribed regions of the genome. The distinction is based on epigenomic features only, which prompted us to identify sequence features that could discriminate Red versus Yellow regions. To this end, we used starcode to cluster 18-mers from both regions.

The total coverage of Red and Yellow chromatin in the *Drosophila* genome is 10.7 and 21.3 Mb respectively, corresponding to 21.5

million and 42.5 million 18-mers respectively. We used starcode to identify clusters of 18-mers within a Levenshtein distance of 3 (the running time was approximately 35 minutes for Red 18-mers and 70 minutes for Yellow 18-mers using 20 cores). The frequencies of the most abundant 18-mers showed that Red chromatin has a lower sequence complexity than Yellow chromatin (not shown). The most discriminating 18-mer was (GA)<sub>9</sub>, which showed a 4.8 fold enrichment in Red versus Yellow chromatin. We counted 832 exact matches and 6,749 inexact matches (over 3,880 distinct 18-mers) in Red chromatin, showing that a larger portion of the signal comes from inexact matches. Taken together, these results show that starcode can be used to cluster short k-mers from metazoan-scale datasets.

## 4 DISCUSSION AND CONCLUSION

Through the parallel poucet search algorithm, starcode implements an exact sequence clustering algorithm that is even faster than widely used heuristics. By design, starcode is tailored to process high throughput sequencing data on multi-core platforms. Our benchmark shows that starcode is superior in speed and accuracy to other methods on the problem of clustering short random sequences. Part of the reason is that this task is stretching the tested software far from their initial design. Sequence clustering algorithms are meant to cluster long biological sequences, while short read mappers are meant to map reads on potentially large genomes. In this respect, starcode fills a need arising from the development of barcoding technologies.

The speed of starcode also makes it proficient at more general clustering problems, such as identifying enriched k-mers in genomes and in experimental data. Here we have given two examples of such applications. In the first, we recover a known enriched 12-mer in the genome of *Neisseria meningitidis*. In the second, we identify a sequence feature that discriminates genomic regions found in two different chromatin types in *Drosophila*. GA repeats are known to be bound in vivo by the protein GAF (van Steensel *et al.*, 2003), which suggests that this transcription factor is instrumental in setting the distinction between Red and Yellow chromatin.

The current version of starcode has been mildly optimized for memory consumption. The memory footprint depends on the number of sequences to cluster (because the sequences of the input set are loaded in memory as a group of tries) and on the mean number of matches per sequence. Every match has to be stored until the clustering phase, which can represent a heavy load. As counterintuitive as it may seem, long queries will usually impose a lower memory footprint because the matches between sequences are less frequent.

The speed of starcode could benefit other applications, such as searching shared subsequences in genome assembly. However, starcode supports only global and not local alignment. A work around is to compare k-mers as we have done in the examples above, but this approach is expected to be memory demanding. Extending the starcode algorithm to support clustering of protein sequences is straightforward but the speed is not expected to meet the level of the algorithm shown here. On the one hand, there are more letters in the protein alphabet, so the shared prefix between consecutive queries is expected to be shorter on average, which defeats the poucet

search. On the other hand, if a threshold distance of the order of 1 per 10 nucleotides is sufficient to identify similar DNA sequences, a distance of 1 per 10 residues is unrealistically low for protein sequences, and this parameter critically determines the speed of starcode. More generally, storing computational intermediates for shared prefixes could find some applications in other algorithms such as short read mappers. The idea of the poucet search seems simple in retrospect, but it is a powerful way to tap into the data structuration provided by string sorting.

## ACKNOWLEDGEMENT

We would like to thank...

**Funding:** The research leading to these results has received funding from the Government of Catalonia (Dept. of Economy and Knowledge) and the Spanish Ministry of Economy and Competitiveness (Centro de Excelencia Severo Ochoa 2013-2017' (SEV-2012-0208). P.C. fellowship is partly financed by the Spanish Ministry of Economy and Competitiveness (State Training Subprogram: predoctoral fellowships for the training of PhD students (FPI) 2013).

## REFERENCES

- Akhtar, W., de Jong, J., Pindyurin, A. V., Pagie, L., Meuleman, W., de Ridder, J., Berns, A., Wessels, L. F. A., van Lohuizen, M., and van Steensel, B. (2013). Chromatin position effects assayed by thousands of reporters integrated in parallel. *Cell*, **154**(4), 914–27.
- Dohm, J. C., Lottaz, C., Borodina, T., and Himmelbauer, H. (2008). Substantial biases in ultra-short read data sets from high-throughput DNA sequencing. *Nucleic Acids Res.*, **36**(16), e105.
- Edgar, R. C. (2010). Search and clustering orders of magnitude faster than BLAST. *Bioinformatics*, **26**(19), 2460–1.
- Eid, J., Fehr, A., Gray, J., Luong, K., Lyle, J., Otto, G., Peluso, P., Rank, D., Baybayan, P., Bettman, B., Bibillo, A., Bjornson, K., Chaudhuri, B., Christians, F., Cicero, R., Clark, S., Dalal, R., Dewinter, A., Dixon, J., Foquet, M., Gaertner, A., Hardenbol, P., Heiner, C., Hester, K., Holden, D., Kearns, G., Kong, X., Kuse, R., Lacroix, Y., Lin, S., Lundquist, P., Ma, C., Marks, P., Maxham, M., Murphy, D., Park, I., Pham, T., Phillips, M., Roy, J., Sebra, R., Shen, G., Sorenson, J., Tomaney, A., Travers, K., Trulson, M., Vieceli, J., Wegener, J., Wu, D., Yang, A., Zaccarin, D., Zhao, P., Zhong, F., Korlach, J., and Turner, S. (2009). Real-time DNA sequencing from single polymerase molecules. *Science*, **323**(5910), 133–8.
- Filion, G. J., van Bommel, J. G., Braunschweig, U., Talhout, W., Kind, J., Ward, L. D., Brugman, W., de Castro, I. J., Kerkhoven, R. M., Bussemaker, H. J., and van Steensel, B. (2010). Systematic protein location mapping reveals five principal chromatin types in *Drosophila* cells. *Cell*, **143**(2), 212–24.
- Fu, L., Niu, B., Zhu, Z., Wu, S., and Li, W. (2012). CD-HIT: accelerated for clustering the next-generation sequencing data. *Bioinformatics*, **28**(23), 3150–2.
- Langmead, B. and Salzberg, S. L. (2012). Fast gapped-read alignment with Bowtie 2. *Nat. Methods*, **9**(4), 357–9.
- Li, H. and Durbin, R. (2009). Fast and accurate short read alignment with Burrows-Wheeler transform. *Bioinformatics*, **25**(14), 1754–60.
- MacKay, D. J. C. (2002). *Information Theory, Inference & Learning Algorithms*. Cambridge University Press, New York, NY, USA.
- Marco-Sola, S., Sammeth, M., Guig, R., and Ribeca, P. (2012). The GEM mapper: fast, accurate and versatile alignment by filtration. *Nat. Methods*, **9**(12), 1185–8.
- Margulies, M., Egholm, M., Altman, W. E., Attiya, S., Bader, J. S., Bembem, L. A., Berka, J., Braverman, M. S., Chen, Y.-J., Chen, Z., Dewell, S. B., Du, L., Fierro, J. M., Gomes, X. V., Godwin, B. C., He, W., Helgesen, S., Ho, C. H., Ho, C. H., Irzyk, G. P., Jando, S. C., Alenquer, M. L. I., Jarvie, T. P., Jirage, K. B., Kim, J.-B., Knight, J. R., Lanza, J. R., Leamon, J. H., Lefkowitz, S. M., Lei, M., Li, J., Lohman, K. L., Lu, H., Makhijani, V. B., McDade, K. E., McKenna, M. P., Myers, E. W., Nickerson, E., Nobile, J. R., Plant, R., Puc, B. P., Ronan, M. T., Roth, G. T., Sarkis, G. J., Simons, J. F., Simpson, J. W., Srinivasan, M., Tartaro, K. R., Tomasz, A., Vogt, K. A., Volkmer, G. A., Wang, S. H., Wang, Y., Weiner, M. P., Yu, P.,

- Begley, R. F., and Rothberg, J. M. (2005). Genome sequencing in microfabricated high-density picolitre reactors. *Nature*, **437**(7057), 376–80.
- Nakamura, K., Oshima, T., Morimoto, T., Ikeda, S., Yoshikawa, H., Shiwa, Y., Ishikawa, S., Linak, M. C., Hirai, A., Takahashi, H., Altaf-Ul-Amin, M., Ogasawara, N., and Kanaya, S. (2011). Sequence-specific error profile of Illumina sequencers. *Nucleic Acids Res.*, **39**(13), e90.
- Needleman, S. B. and Wunsch, C. D. (1970). A general method applicable to the search for similarities in the amino acid sequence of two proteins. *J. Mol. Biol.*, **48**(3), 443–53.
- Schepers, K., Swart, E., van Heijst, J. W. J., Gerlach, C., Castrucci, M., Sie, D., Heimerikx, M., Velds, A., Kerkhoven, R. M., Arens, R., and Schumacher, T. N. M. (2008). Dissecting T cell lineage relationships by cellular barcoding. *J. Exp. Med.*, **205**(10), 2309–18.
- Smith, H. O., Gwinn, M. L., and Salzberg, S. L. (1999). DNA uptake signal sequences in naturally transformable bacteria. *Res. Microbiol.*, **150**(9-10), 603–16.
- Ukkonen, E. (1995). On-line construction of suffix trees. *Algorithmica*, **14**(3), 249–260.
- van Steensel, B., Delrow, J., and Bussemaker, H. J. (2003). Genomewide analysis of Drosophila GAGA factor target genes reveals context-dependent DNA binding. *Proc. Natl. Acad. Sci. U.S.A.*, **100**(5), 2580–5.